

# Bellman Equations

---

## Reward

---

$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

## Policy

---

Policy is a function  $\Pi(s, a)$  of the state and the current action. It returns the probability of taking action  $a$  in state  $s$ .

## State value function

---

$$V^{\pi}(s) = E_{\pi}[R_t | s_t = s]$$

It is the expected return when starting from state  $s$  according to policy  $\pi$

## Action value function

---

$$Q^{\pi}(s, a) = E_{\pi}[R_t | s_t = s, a_t = a]$$

It is the expected return given  $s$  and  $a$  under  $\pi$

## Bellman equation for state value function

---

$$V^{\pi}(s) = \sum_a \pi(s, a) \sum_{s'} P_{ss'}^a [R_{ss'}^a + \gamma V^{\pi}(s')]$$

## Bellman equation for action value function

---

$$Q^{\pi}(s, a) = \sum_{s'} P_{ss'}^a [R_{ss'}^a + \gamma \sum_{a'} \pi(s', a') Q^{\pi}(s', a')]$$

# Training algorithms

---

## Q-learning

---

### Deterministic Bellman

$$Q(s, a) = r + \gamma * \max(Q(s', a'))$$

```
# Randomize current QValues
rand_qvals = Q[state] + torch.rand(1, number_of_actions)/1000
# Get an action given the current QValues
action = torch.max(rand_qvals, 1)[1][0].item()
# Produce a new state given the chosen action
new_state, reward, done, info = env.step(action)
# Update QValues given current reward and QValue
Q[state, action] = reward + gamma * torch.max(Q[new_state])
state = new_state
```

### Stochastic Q-Learning

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha[r + \gamma * \max(Q(s', a'))]$$

```
rand_qvals = Q[state] + torch.rand(1, number_of_actions)/1000
action = torch.max(rand_qvals, 1)[1][0].item()
new_state, reward, done, info = env.step(action)
Q[state, action] = (1 - learning_rate) * Q[state, action] + learning_rate *
(reward + gamma * torch.max(Q[new_state]))
state = new_state
```

### E-Greedy decay

---

```

random_for_egreedy = torch.rand(1)[0].item()
if random_for_egreedy > egreedy:
    rand_qvals = Q[state] + torch.rand(1, number_of_actions)/1000
    action = torch.max(rand_qvals, 1)[1][0].item()
else:
    action = env.action_space.sample()

if egreedy > egreedy_final:
    egreedy *= egreedy_decay

new_state, reward, done, info = env.step(action)
Q[state, action] = reward + gamma * torch.max(Q[new_state])
state = new_state

```

## Gradient descent - Neural Network

$$Q(s, a) = f(s, a)$$

$f(.)$  is approximated using a neural network

## Standard optimization

```

def optimize(self, state, action, new_state, reward, done):
    state = torch.Tensor(state).to(device)
    new_state = torch.Tensor(new_state).to(device)
    reward = torch.Tensor([reward]).to(device)

    if done:
        target_value = reward
    else:
        new_state_values = self.nn(new_state).detach()
        max_new_state_values = torch.max(new_state_values)
        target_value = reward + gamma * max_new_state_values

    predicted_value = self.nn(state)[action]
    loss = self.loss_func(predicted_value, target_value)
    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()

```

## Experience replay

```
def optimize(self):

    if len(memory) < batch_size:
        return

    state, action, new_state, reward, done = memory.sample(batch_size)
    state = torch.Tensor(state).to(device)
    new_state = torch.Tensor(new_state).to(device)
    reward = torch.Tensor(reward).to(device)
    action = torch.LongTensor(action).to(device)
    done = torch.Tensor(done).to(device)

    new_state_values = self.nn(new_state).detach()
    max_new_state_values = torch.max(new_state_values, 1)[0]
    target_value = reward + (1 - done) * gamma * max_new_state_values

    predicted_value = self.nn(state).gather(1, action.unsqueeze(1)).squeeze(1)
    loss = self.loss_func(predicted_value, target_value)
    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()
```

## Target Net

```
def optimize(self):

    if len(memory) < batch_size:
        return

    state, action, new_state, reward, done = memory.sample(batch_size)
    state = torch.Tensor(state).to(device)
    new_state = torch.Tensor(new_state).to(device)
    reward = torch.Tensor(reward).to(device)
    action = torch.LongTensor(action).to(device)
    done = torch.Tensor(done).to(device)

    new_state_values = self.target_nn(new_state).detach()
    max_new_state_values = torch.max(new_state_values, 1)[0]
    target_value = reward + (1 - done) * gamma * max_new_state_values

    predicted_value = self.nn(state).gather(1, action.unsqueeze(1)).squeeze(1)
```

```
loss = self.loss_func(predicted_value, target_value)
self.optimizer.zero_grad()
loss.backward()
if clip_error:
    for param in self.nn.parameters():
        param.grad.data.clamp_(-1,1)
        self.optimizer.step()

    if self.update_target_counter % update_target_frequency == 0:
        self.target_nn.load_state_dict(self.nn.state_dict())

        self.update_target_counter + 1
```